

Chapters 5: Propositions and Inference

DIT410/TIN172 Artificial Intelligence

Peter Ljunglöf

modified from slides by Poole & Mackworth

(Licensed under Creative Commons BY-NC-SA v4.0)

1 April, 2015

For when I am presented with a false theorem, I do not need to examine or even to know the demonstration, since I shall discover its falsity *a posteriori* by means of an easy experiment, that is, by a calculation, costing no more than paper and ink, which will show the error no matter how small it is. . .

And if someone would doubt my results, I should say to him: "Let us calculate, Sir," and thus by taking to pen and ink, we should soon settle the question.

—Gottfried Wilhelm Leibniz [1677]

Outline

- ➊ *Propositions (5.1–5.2)*
- ➋ *Proofs (5.2.2)*
- ➌ *Bottom-up proof procedure (5.2.2.1)*
- ➍ *Top-down proof procedure (5.2.2.2)*
- ➎ *Complete knowledge assumption (5.5)*

Outline

- ➊ *Propositions (5.1–5.2)*
- ➋ *Proofs (5.2.2)*
- ➌ *Bottom-up proof procedure (5.2.2.1)*
- ➍ *Top-down proof procedure (5.2.2.2)*
- ➎ *Complete knowledge assumption (5.5)*

Interpretations, models and propositions

- An *interpretation* is an assignment of values to variables.
- A *model* is an interpretation that satisfies the constraints.

Often we don't want to just find a model, but we want to know what is true in all models:

- A *proposition* is statement that is true or false in each interpretation.
- Propositions are built using *logical connectives*.

Syntax of propositional calculus

Propositions are built from simpler propositions using logical connectives. A proposition is either:

- an atomic proposition (also called atom, symbol or boolean variable),
- or a compound proposition of the form:

$\neg p$: the negation of p

$p \wedge q$: the conjunction of p and q

$p \vee q$: the disjunction of p and q

$p \rightarrow q$: the implication of q from p

$p \leftrightarrow q$: the equivalence of p and q

The precedence of the connectives is in the above order.

Semantics of propositional calculus

An interpretation is a function π that maps atoms to $\{true, false\}$:

- if $\pi(a) = true$ (*false*), we say that atom a is true (*false*) in the interpretation
- we can also think of π as the set of atoms that map to true

The interpretation maps each proposition to a *truth value*.

The value of a compound proposition is built using this *truth table*:

p	q	$\neg q$	$p \wedge q$	$p \vee q$	$p \rightarrow q$	$p \leftrightarrow q$
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>
<i>true</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>
<i>false</i>	<i>true</i>		<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>
<i>false</i>	<i>false</i>		<i>false</i>	<i>false</i>	<i>true</i>	<i>true</i>

Simple language: Propositional definite clauses

The *definite clauses* are a subset of all propositions:

- An **atom** is a symbol starting with a lower case letter
- A **body** is of the form $b_1 \wedge \dots \wedge b_k$ where $b_1 \dots b_k$ are atoms
 - ▶ k can be 1, in which case the body is a single atom
- A **definite clause** is an atom or is an implication of the form $b \rightarrow h$ where b is a body and h is an atom
 - ▶ usually we write the clause backwards, $h \leftarrow b$
- A **knowledge base** is a set of definite clauses

Definite Clauses

Which of the following are definite clauses?

- (a) $happy \leftarrow sad$
- (b) $blimsy$
- (c) $old \wedge wise \leftarrow teenager$
- (d) $happy \wedge sad$
- (e) $glad \leftarrow happy \wedge sad$
- (f) $green \vee blue \leftarrow \neg red$
- (g) $glad \leftarrow happy \wedge sad \wedge mad \wedge bad$
- (h) $glad \leftarrow happy \wedge rad \leftarrow sad \wedge mad \wedge bad$
- (i) $happy \leftarrow happy$

Human's vs computer's view of semantics

In the user's mind:

- *light1_broken*: light #1 is broken
- *sw_up*: switch is up
- *power*: there is power in the building
- *unlit_light1*: light #1 isn't lit
- *lit_light2*: light #2 is lit

In the computer:

$light1_broken \leftarrow sw_up$
 $\wedge power \wedge unlit_light1.$
 $sw_up.$
 $power \leftarrow lit_light2.$
 $unlit_light1.$
 $lit_light2.$

Logical conclusion: *light1_broken*

- The computer doesn't know the meaning of the symbols
- The user can interpret the symbol using their meaning

Semantics for definite clauses

Simplified semantics for definite clauses:

- An **interpretation** I assigns a truth value to each atom.
- A rule $h \leftarrow b_1 \wedge \dots \wedge b_k$ is *false in I* if:
 - ▶ h is *false in I* , and
 - ▶ all b_i 's are *true in I* .
- Otherwise the rule is *true in I* .
- A knowledge base KB is *true in I* if and only if every clause in KB is *true in I* .

Models and logical consequence

- A **model** of a set of clauses is an interpretation in which all the clauses are *true*.
- Assuming that KB is a set of clauses and g is a body:
 - ▶ g is a **logical consequence** of KB , written $KB \models g$, if g is *true* in every model of KB .
- That is, $KB \models g$ if there is no interpretation in which KB is *true* and g is *false*.

Simple example

$$KB = \begin{cases} p \leftarrow q. \\ q. \\ r \leftarrow s. \end{cases}$$

	<i>p</i>	<i>q</i>	<i>r</i>	<i>s</i>	model?
<i>I</i> ₁	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>	
<i>I</i> ₂	<i>false</i>	<i>false</i>	<i>false</i>	<i>false</i>	
<i>I</i> ₃	<i>true</i>	<i>true</i>	<i>false</i>	<i>false</i>	
<i>I</i> ₄	<i>true</i>	<i>true</i>	<i>true</i>	<i>false</i>	
<i>I</i> ₅	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	

Which of *p*, *q*, *r*, *s*, *t* logically follow from KB?

Simple example

$$KB = \begin{cases} p \leftarrow q. \\ q. \\ r \leftarrow s. \end{cases}$$

	<i>p</i>	<i>q</i>	<i>r</i>	<i>s</i>	model?
<i>I</i> ₁	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>	is a model of <i>KB</i>
<i>I</i> ₂	<i>false</i>	<i>false</i>	<i>false</i>	<i>false</i>	not a model of <i>KB</i>
<i>I</i> ₃	<i>true</i>	<i>true</i>	<i>false</i>	<i>false</i>	is a model of <i>KB</i>
<i>I</i> ₄	<i>true</i>	<i>true</i>	<i>true</i>	<i>false</i>	is a model of <i>KB</i>
<i>I</i> ₅	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	not a model of <i>KB</i>

Which of *p*, *q*, *r*, *s*, *t* logically follow from *KB*?

$KB \models p$, $KB \models q$, $KB \not\models r$, $KB \not\models s$, $KB \not\models t$

User's view of semantics

- Step 1* Choose a task domain: **intended interpretation.**
- Step 2* Associate an atom with each proposition you want to represent.
- Step 3* Tell the system clauses that are true in the intended interpretation: **axiomatising the domain.**
- Step 4* Ask questions about the intended interpretation.
- If $KB \models g$, then g must be true in the intended interpretation.
- Step 5* Users can interpret the answer using their intended interpretation of the symbols.

Computer's view of semantics

- The computer doesn't have access to the intended interpretation.
 - ▶ All it knows is the knowledge base KB .
- It can determine if a formula is a logical consequence of KB .
- If $KB \models g$ then g must be true in the intended interpretation.
- If $KB \not\models g$ then there is a model of KB in which g is false.


Computer's view of semantics

- The computer doesn't have access to the intended interpretation.
 - ▶ All it knows is the knowledge base KB .
- It can determine if a formula is a logical consequence of KB .
- If $KB \models g$ then g must be true in the intended interpretation.
- If $KB \not\models g$ then there is a model of KB in which g is false.
 - ▶ This could be the intended interpretation!

Computer's view of semantics

- The computer doesn't have access to the intended interpretation.
 - ▶ All it knows is the knowledge base KB .
- It can determine if a formula is a logical consequence of KB .
- If $KB \models g$ then g must be true in the intended interpretation.
- If $KB \not\models g$ then there is a model of KB in which g is false.
 - ▶ This could be the intended interpretation!
 - ▶ (then there is something wrong with the axiomatisation)

Example: Electrical environment



Representing the electrical environment

*light*_{l₁}.

*light*_{l₂}.

*down*_{s₁}.

*up*_{s₂}.

*up*_{s₃}.

*ok*_{l₁}.

*ok*_{l₂}.

*ok*_{cb₁}.

*ok*_{cb₂}.

*live*_{outside}.

*lit*_{l₁} \leftarrow *live*_{w₀} \wedge *ok*_{l₁}

*live*_{w₀} \leftarrow *live*_{w₁} \wedge *up*_{s₂}.

*live*_{w₀} \leftarrow *live*_{w₂} \wedge *down*_{s₂}.

*live*_{w₁} \leftarrow *live*_{w₃} \wedge *up*_{s₁}.

*live*_{w₂} \leftarrow *live*_{w₃} \wedge *down*_{s₁}.

*lit*_{l₂} \leftarrow *live*_{w₄} \wedge *ok*_{l₂}.

*live*_{w₄} \leftarrow *live*_{w₃} \wedge *up*_{s₃}.

*live*_{p₁} \leftarrow *live*_{w₃}.

*live*_{w₃} \leftarrow *live*_{w₅} \wedge *ok*_{cb₁}.

*live*_{p₂} \leftarrow *live*_{w₆}.

*live*_{w₆} \leftarrow *live*_{w₅} \wedge *ok*_{cb₂}.

*live*_{w₅} \leftarrow *live*_{outside}.

Outline

- 1 *Propositions (5.1–5.2)*
- 2 *Proofs (5.2.2)*
- 3 *Bottom-up proof procedure (5.2.2.1)*
- 4 *Top-down proof procedure (5.2.2.2)*
- 5 *Complete knowledge assumption (5.5)*

Proofs

- A **proof** is a mechanically derivable demonstration that a formula logically follows from a knowledge base.
- $KB \vdash g$ means that g can be derived from knowledge base KB , using a proof procedure.
- Recall that $KB \models g$ means g is true in all models of KB .
- A proof procedure is **sound** if $KB \vdash g$ implies $KB \models g$.
- A proof procedure is **complete** if $KB \models g$ implies $KB \vdash g$.

Outline

- 1 *Propositions (5.1–5.2)*
- 2 *Proofs (5.2.2)*
- 3 *Bottom-up proof procedure (5.2.2.1)*
- 4 *Top-down proof procedure (5.2.2.2)*
- 5 *Complete knowledge assumption (5.5)*

Bottom-up proof procedure

- There is one single **rule of derivation**, a generalized form of *modus ponens*:

If “ $h \leftarrow b_1 \wedge \dots \wedge b_m$ ” is a clause in the knowledge base, and each b_i has been derived, then h can be derived.

(This rule also covers the case when $m = 0$)

- This is called **forward chaining** on the clause.

Bottom-up proof procedure

After the following procedure, $KB \vdash g$ iff $g \in C$:

$C := \{\}$

repeat

select clause " $h \leftarrow b_1 \wedge \dots \wedge b_m$ " in KB such that

$b_i \in C$ for all i , and $h \notin C$

$C := C \cup \{h\}$

until no more clauses can be selected

Example

$$a \leftarrow b \wedge c.$$

$$b \leftarrow d \wedge e.$$

$$b \leftarrow g \wedge e.$$

$$c \leftarrow e.$$

$$d.$$

$$e.$$

$$f \leftarrow a \wedge g.$$

Soundness of bottom-up proof procedure

Proof by contradiction:

- Assume there is a g such that $KB \vdash g$ and $KB \not\models g$.
- Then there must be a *first* atom added to C that isn't true in every model of KB . Call it h .
 - ▶ Let's call this model I , in which h isn't *true*.
- There must be a clause in KB of form:

$$h \leftarrow b_1 \wedge \dots \wedge b_m$$

Each b_i is true in I , but h is false in I . So this clause is false in I . Therefore I isn't a model of KB .

- Contradiction.

Therefore, **if $KB \vdash g$ then $KB \models g$.**

Fixed point

- The C generated at the end of the bottom-up algorithm is called a **fixed point**.
- Let I be the interpretation in which every element of the fixed point is true and every other atom is false.

Theorem: I is a model of KB .

Proof: Suppose $h \leftarrow b_1 \wedge \dots \wedge b_m$ in KB is false in I .

Then h is false and each b_i is true in I .

Thus h can be added to C .

This is a contradiction to C being the fixed point.

- I is called a **minimal model**.

Completeness of the bottom-up proof procedure

- Assume that $KB \models g$.
- Then g is true in all models of KB .
- Thus g is true in the minimal model.
- Thus g is in the fixed point.
- Thus g is generated by the bottom up algorithm.
- Thus $KB \vdash g$.

Therefore, if $KB \models g$ then $KB \vdash g$.

Outline

- 1 *Propositions (5.1–5.2)*
- 2 *Proofs (5.2.2)*
- 3 *Bottom-up proof procedure (5.2.2.1)*
- 4 *Top-down proof procedure (5.2.2.2)*
- 5 *Complete knowledge assumption (5.5)*

Top-down proof procedure

A **query** is a body that we want to determine whether it is a logical consequence of KB .

Idea: search backward from the query.

- An **answer clause** is of the form:

$$yes \leftarrow a_1 \wedge \cdots \wedge a_i \wedge \cdots \wedge a_m$$

- **SLD resolution** of this answer clause on atom a_i with the clause:

$$a_i \leftarrow b_1 \wedge \cdots \wedge b_p$$

results in the answer clause:

$$yes \leftarrow a_1 \wedge \cdots \wedge a_{i-1} \wedge b_1 \wedge \cdots \wedge b_p \wedge a_{i+1} \wedge \cdots \wedge a_m$$

Derivations

- An **answer** is an answer clause with $m = 0$.
 - ▶ i.e., it is the answer clause “ $yes \leftarrow$ ”
- A **derivation** of query “ $?q_1 \wedge \dots \wedge q_k$ ” from KB is a sequence of answer clauses $\gamma_0, \gamma_1, \dots, \gamma_n$ such that
 - ▶ γ_0 is the answer clause “ $yes \leftarrow q_1 \wedge \dots \wedge q_k$ ”,
 - ▶ γ_i is obtained by resolving γ_{i-1} with a clause in KB , and
 - ▶ γ_n is an answer.

Top-down definite clause interpreter

To solve the query $?q_1 \wedge \dots \wedge q_k$:

$ac := \text{"yes"} \leftarrow q_1 \wedge \dots \wedge q_k$

repeat

select atom a_i from the body of ac

choose clause C from KB with a_i as head

replace a_i in the body of ac by the body of C

until ac is an answer

Nondeterministic choice

The top-down interpreter uses two kinds of nondeterministic choice:

don't-care If one selection doesn't lead to a solution,
there is no point trying other alternatives.

- “select atom a_i from the body of ac ”

don't-know If one choice doesn't lead to a solution,
there might be other choices that will.

- “choose clause C from KB with a_i as head”

Example: Successful derivation

$$a \leftarrow b \wedge c$$

$$c \leftarrow e$$

$$f \leftarrow j \wedge e$$

$$a \leftarrow e \wedge f$$

$$d \leftarrow k$$

$$f \leftarrow c$$

$$b \leftarrow f \wedge k$$

$$e$$

$$j \leftarrow c$$

Query: $?a$

Example: Successful derivation

$$\begin{array}{lll}
 a \leftarrow b \wedge c & a \leftarrow e \wedge f & b \leftarrow f \wedge k \\
 c \leftarrow e & d \leftarrow k & e \\
 f \leftarrow j \wedge e & f \leftarrow c & j \leftarrow c
 \end{array}$$

Query: $?a$

$$\begin{array}{ll}
 \gamma_0 : \text{yes} \leftarrow a & (a \leftarrow e \wedge f) \\
 \gamma_1 : \text{yes} \leftarrow e \wedge f & (e) \\
 \gamma_2 : \text{yes} \leftarrow f & (f \leftarrow c) \\
 \gamma_3 : \text{yes} \leftarrow c & (c \leftarrow e) \\
 \gamma_4 : \text{yes} \leftarrow e & (e) \\
 \gamma_5 : \text{yes} \leftarrow &
 \end{array}$$

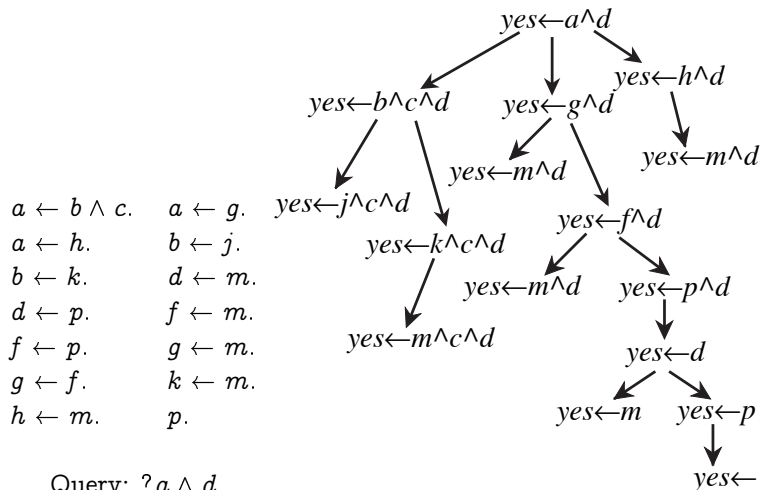
Example: Failing derivation

$$\begin{array}{lll}
 a \leftarrow b \wedge c & a \leftarrow e \wedge f & b \leftarrow f \wedge k \\
 c \leftarrow e & d \leftarrow k & e \\
 f \leftarrow j \wedge e & f \leftarrow c & j \leftarrow c
 \end{array}$$

Query: $?a$

$$\begin{array}{ll}
 \gamma_0 : \text{yes} \leftarrow a & (a \leftarrow b \wedge c) \\
 \gamma_1 : \text{yes} \leftarrow b \wedge c & (b \leftarrow f \wedge k) \\
 \gamma_2 : \text{yes} \leftarrow f \wedge k \wedge c & (f \leftarrow c) \\
 \gamma_3 : \text{yes} \leftarrow c \wedge k \wedge c & (c \leftarrow e) \\
 \gamma_4 : \text{yes} \leftarrow e \wedge k \wedge c & (e) \\
 \gamma_5 : \text{yes} \leftarrow k \wedge c & \text{fail}
 \end{array}$$

Search graph for SLD resolution



Outline

- 1 *Propositions (5.1–5.2)*
- 2 *Proofs (5.2.2)*
- 3 *Bottom-up proof procedure (5.2.2.1)*
- 4 *Top-down proof procedure (5.2.2.2)*
- 5 *Complete knowledge assumption (5.5)*

Complete knowledge assumption

- Often you want to assume that your knowledge is complete.
 - ▶ *Example:* you can state what switches are up and then the agent can assume that the other switches are down.
 - ▶ *Example:* assume that a database is complete, e.g., of what students are enrolled in a course.
- The definite clause language is **monotonic**:
 - ▶ adding clauses can't invalidate a previous conclusion.
- Under the *complete knowledge assumption*, the system is **non-monotonic**:
 - ▶ adding clauses can invalidate a previous conclusion.

Completion of a knowledge base

- Suppose the rules for atom a are,

$$a \leftarrow b_1 \quad \dots \quad a \leftarrow b_n$$

or equivalently $a \leftarrow b_1 \vee \dots \vee b_n$.

- The *complete knowledge assumption* (CKA) says that if a is true, one of the b_i must be true:

$$a \rightarrow b_1 \vee \dots \vee b_n$$

- Under the CKA, the meaning of the clauses for a are:

$$a \leftrightarrow b_1 \vee \dots \vee b_n$$

which is called **Clark's completion.**

Clark's completion of a KB

- Clark's completion of a knowledge base consists of the completion of every atom.
- If you have an atom a with no clauses, the completion is $a \leftrightarrow \text{false}$.
- You can interpret negations in the body of clauses:
 - ▶ $\sim a$ means that a is false under the CKA
 - ▶ this is **negation as failure**

Bottom-up negation-as-failure interpreter

```

 $C := \{\}$ 
repeat
  either
    select  $h$  such that there is a rule " $h \leftarrow b_1 \wedge \dots \wedge b_m$ "  $\in KB$ 
      where every  $b_i \in C$ , and  $h \notin C$ 
     $C := C \cup \{h\}$ 
  or
    select  $h$  such that for every rule " $h \leftarrow b_1 \wedge \dots \wedge b_m$ "  $\in KB$ 
      either  $\sim b_i \in C$  for some  $b_i$ 
      or  $g \in C$  for some  $b_i = \sim g$ 
     $C := C \cup \{\sim h\}$ 
until no more selections are possible

```

Negation-as-failure example

$$p \leftarrow q \wedge \sim r.$$

$$p \leftarrow s.$$

$$q \leftarrow \sim s.$$

$$r \leftarrow \sim t.$$

$$t.$$

$$s \leftarrow w.$$

Top-down negation-as-failure proof procedure

If the proof for a fails, you can conclude $\sim a$.

Failure can be defined recursively:

- Suppose you have rules for atom a :

$$a \leftarrow b_1 \quad \dots \quad a \leftarrow b_n$$

- If each body b_i fails, a fails.
- A body fails if one of the conjuncts in the body fails.
- Note that you need *finite* failure.
 - ▶ counter-example: $p \leftarrow p$

Default reasoning

Exceptions are very common in ontological databases:

- Birds fly.
- Emus and tiny birds don't.
- Hummingbirds are tiny birds that can fly.

Negation-as-failure can be used for *default reasoning*:

$$\begin{aligned}
 \textit{flies} &\leftarrow \textit{bird} \wedge \sim \textit{ab}_{\textit{flying}} \\
 \textit{ab}_{\textit{flying}} &\leftarrow \textit{emu} \wedge \sim \textit{ab}_{\textit{emu}} \\
 \textit{ab}_{\textit{flying}} &\leftarrow \textit{tiny} \wedge \sim \textit{ab}_{\textit{tiny}} \\
 \textit{ab}_{\textit{tiny}} &\leftarrow \textit{hummingbird} \wedge \sim \textit{ab}_{\textit{hummingbird}}
 \end{aligned}$$