

# CHAPTERS 3–4: MORE SEARCH ALGORITHMS

DIT410/TIN174, Artificial Intelligence

Peter Ljunglöf

28 March, 2017

# TABLE OF CONTENTS

## Heuristic search (R&N 3.5–3.6)

- Greedy best-first search (3.5.1)
- A\* search (3.5.2)
- Admissible and consistent heuristics (3.6–3.6.2)

## More search strategies (R&N 3.4–3.5)

- Iterative deepening (3.4.4–3.4.5)
- Bidirectional search (3.4.6)
- Memory-bounded A\* (3.5.3)

## Local search (R&N 4.1)

- Hill climbing search (4.1.1–4.1.2)
- Population-based methods (4.1.3–4.1.4)
- Evaluating randomized algorithms (not in R&N)

# HEURISTIC SEARCH (R&N 3.5–3.6)

GREEDY BEST-FIRST SEARCH (3.5.1)

A\* SEARCH (3.5.2)

ADMISSIBLE AND CONSISTENT HEURISTICS (3.6–3.6.2)

# THE GENERIC TREE SEARCH ALGORITHM

*Tree search*: Don't check if nodes are visited multiple times

```
function Search(graph, initialState, goalState):  
  initialise frontier using the initialState  
  while frontier is not empty:  
    select and remove node from frontier  
    if node.state is a goalState then return node  
    for each child in ExpandChildNodes(node, graph):  
      add child to frontier  
  return failure
```

# DEPTH-FIRST AND BREADTH-FIRST SEARCH

## THESE ARE THE TWO BASIC SEARCH ALGORITHMS

### Depth-first search (DFS)

- implement the frontier as a Stack
- space complexity:  $O(bm)$
- incomplete: might fall into an infinite loop, doesn't return optimal solution

### Breadth-first search (BFS)

- implement the frontier as a Queue
- space complexity:  $O(b^m)$
- complete: always finds a solution, if there is one
- (when edge costs are constant, BFS is also optimal)

# COST-BASED SEARCH

IMPLEMENT THE FRONTIER AS A PRIORITY QUEUE, ORDERED BY  $f(n)$

Uniform-cost search (this is not a heuristic algorithm)

- expand the node with the lowest path cost
- $f(n) = g(n)$
- complete and optimal

Greedy best-first search

- expand the node which is closest to the goal (according to some heuristics)
- $f(n) = h(n)$
- incomplete: might fall into an infinite loop, doesn't return optimal solution

A\* search

- expand the node which has the lowest estimated cost from start to goal
- $f(n) = g(n) + h(n)$  = estimated cost of the cheapest solution through  $n$
- complete and optimal (if  $h(n)$  is admissible/consistent)

## A\* TREE SEARCH IS OPTIMAL!

A\* always finds an optimal solution first, provided that:

- the branching factor is finite,
- arc costs are *bounded above zero*  
(i.e., there is some  $\epsilon > 0$  such that all of the arc costs are greater than  $\epsilon$ ), and
- $h(n)$  is *admissible*
- i.e.,  $h(n)$  is *nonnegative* and an *underestimate* of the cost of the shortest path from  $n$  to a goal node.

# THE GENERIC GRAPH SEARCH ALGORITHM


*Tree search*: Don't check if nodes are visited multiple times

*Graph search*: Keep track of visited nodes

```
function Search(graph, initialState, goalState):  
    initialise frontier using the initialState  
    initialise exploredSet to the empty set  
    while frontier is not empty:  
        select and remove node from frontier  
        if node.state is a goalState then return node  
        add node to exploredSet  
        for each child in ExpandChildNodes(node, graph):  
            if child is not in frontier or exploredSet:  
                add child to frontier  
    return failure
```



## GRAPH-SEARCH = MULTIPLE-PATH PRUNING




Graph search keeps track of visited nodes, so we don't visit the same node twice.

- Suppose that the first time we visit a node is not via the most optimal path  
⇒ then graph search will return a suboptimal path
- Under which circumstances can we guarantee that A\* graph search is optimal?

# WHEN IS A\* GRAPH SEARCH OPTIMAL?

If  $h$  is *consistent*, then A\* graph search is optimal:

- Consistency is defined as:  $h(n') \leq \text{cost}(n', n) + h(n)$  for all arcs  $(n', n)$
- **Lemma:** the  $f$  values along any path  $[\dots, n', n, \dots]$  are nondecreasing:
  - **Proof:**  $g(n) = g(n') + \text{cost}(n', n)$ , therefore:
  - $f(n) = g(n) + h(n) = g(n') + \text{cost}(n', n) + h(n) \geq g(n') + h(n')$ ;
  - therefore:  $f(n) \geq f(n')$ , i.e.,  $f$  is nondecreasing
- **Theorem:** whenever A\* expands a node  $n$ , the optimal path to  $n$  has been found
  - **Proof:** Assume this is not true;
  - then there must be some  $n'$  still on the frontier, which is on the optimal path to  $n$ ;
  - but  $f(n') \leq f(n)$ ;
  - and then  $n'$  must already have been expanded  $\implies$  *contradiction!*



## STATE-SPACE CONTOURS

The  $f$  values in  $A^*$  are nondecreasing, therefore:

**first**  $A^*$  expands all nodes with  $f(n) < C$

**then**  $A^*$  expands all nodes with  $f(n) = C$

**finally**  $A^*$  expands all nodes with  $f(n) > C$

$A^*$  will not expand any nodes with  $f(n) > C^*$ ,  
where  $C^*$  is the cost of an optimal solution.

## SUMMARY OF OPTIMALITY OF A\*

A\* *tree search* is optimal if:

- the heuristic function  $h(n)$  is **admissible**
- i.e.,  $h(n)$  is nonnegative and an underestimate of the actual cost
- i.e.,  $h(n) \leq \text{cost}(n, \text{goal})$ , for all nodes  $n$

A\* *graph search* is optimal if:

- the heuristic function  $h(n)$  is **consistent** (or monotone)
- i.e.,  $|h(m) - h(n)| \leq \text{cost}(m, n)$ , for all arcs  $(m, n)$

## SUMMARY OF TREE SEARCH STRATEGIES

Search strategy	Frontier selection	Halts if solution?	Halts if no solution?	Space usage
Depth first	Last node added	<i>No</i>	<i>No</i>	<i>Linear</i>
Breadth first	First node added	<i>Yes</i>	<i>No</i>	<i>Exp</i>
Greedy best first	Minimal $h(n)$	<i>No</i>	<i>No</i>	<i>Exp</i>
Uniform cost	Minimal $g(n)$	<i>Optimal</i>	<i>No</i>	<i>Exp</i>
A*	$f(n) = g(n) + h(n)$	<i>Optimal*</i>	<i>No</i>	<i>Exp</i>

*\*Provided that  $h(n)$  is admissible.*

**Halts if:** If there is a path to a goal, it can find one, even on infinite graphs.

**Halts if no:** Even if there is no solution, it will halt on a finite graph (with cycles).


**Space:** Space complexity as a function of the length of the current path.

## RECAPITULATION: HEURISTICS FOR THE 8 PUZZLE

$h_1(n)$  = number of misplaced tiles

$h_2(n)$  = total Manhattan distance

(i.e., no. of squares from desired location of each tile)



$$h_1(\text{StartState}) = 8$$

$$h_2(\text{StartState}) = 3+1+2+2+2+3+3+2 = 18$$

# DOMINATING HEURISTICS

If (admissible)  $h_2(n) \geq h_1(n)$  for all  $n$ ,  
then  $h_2$  **dominates**  $h_1$  and is better for search.

Typical search costs (for 8-puzzle):

<b>depth = 14</b>	DFS $\approx$ 3,000,000 nodes $A^*(h_1) = 539$ nodes $A^*(h_2) = 113$ nodes
-------------------	---

---

<b>depth = 24</b>	DFS $\approx$ 54,000,000,000 nodes $A^*(h_1) = 39,135$ nodes $A^*(h_2) = 1,641$ nodes
-------------------	---

Given any admissible heuristics  $h_a, h_b$ , the **maximum** heuristics  $h(n)$  is also admissible and dominates both:

$$h(n) = \max(h_a(n), h_b(n))$$

## HEURISTICS FROM A RELAXED PROBLEM

Admissible heuristics can be derived from the exact solution cost of a relaxed problem:

- If the rules of the 8-puzzle are relaxed so that a tile can move anywhere, then  $h_1(n)$  gives the shortest solution
- If the rules are relaxed so that a tile can move to any adjacent square, then  $h_2(n)$  gives the shortest solution

**Key point:** the optimal solution cost of a relaxed problem is never greater than the optimal solution cost of the real problem



# NON-ADMISSIBLE (NON-CONSISTENT) A\* SEARCH

A\* search with admissible (consistent) heuristics is optimal

But what happens if the heuristics is non-admissible?

- i.e., what if  $h(n) > c(n, goal)$ , for some  $n$ ?
- the solution is not guaranteed to be optimal...
- ...but it will find *some* solution!

Why would we want to use a non-admissible heuristics?

- sometimes it's easier to come up with a heuristics that is almost admissible
- and, often, the search terminates faster!

## EXAMPLE DEMO

Here is an example demo of several different search algorithms, including A\*. Furthermore you can play with different heuristics:

<http://qiao.github.io/PathFinding.js/visual/>

Note that this demo is tailor-made for planar grids, which is a special case of all possible search graphs.

# **MORE SEARCH STRATEGIES (R&N 3.4–3.5)**

**ITERATIVE DEEPENING (3.4.4–3.4.5)**

**BIDIRECTIONAL SEARCH (3.4.6)**

**MEMORY-BOUNDED HEURISTIC SEARCH (3.5.3)**

# ITERATIVE DEEPENING

BFS is guaranteed to halt but uses exponential space.

DFS uses linear space, but is not guaranteed to halt.


*Idea*: take the best from BFS and DFS — recompute elements of the frontier rather than saving them.

- Look for paths of depth 0, then 1, then 2, then 3, etc.
- Depth-bounded DFS can do this in linear space.

**Iterative deepening search** calls depth-bounded DFS with increasing bounds:

- If a path cannot be found at *depth-bound*, look for a path at *depth-bound* + 1.
- Increase *depth-bound* when the search fails unnaturally (i.e., if *depth-bound* was reached).

# ITERATIVE DEEPENING EXAMPLE



Depth bound = 3

# ITERATIVE-DEEPENING SEARCH

```
function IDSearch(graph, initialState, goalState)  
  for limit in 0, 1, 2, ...:  
    result := DepthLimitedSearch([initialState], limit)  
    if result ≠ cutoff then return result  
  
function DepthLimitedSearch([n0, ... , nk], limit):  
  if nk is a goalState then return path [n0, ... , nk]  
  else if limit = 0 then return cutoff  
  else:  
    failureType := failure  
    for each neighbor n of nk:  
      result := DepthLimitedSearch([n0, ... , nk, n], limit−1)  
      if result is a path then return result  
      else if result = cutoff then failureType := cutoff  
  return failureType
```

# ITERATIVE DEEPENING COMPLEXITY

Complexity with solution at depth  $k$  and branching factor  $b$ :

level	breadth-first	iterative deepening	# nodes
1	1	$k$	$b$
2	1	$k - 1$	$b^2$
$\vdots$	$\vdots$	$\vdots$	$\vdots$
$k - 1$	1	2	$b^{k-1}$
$k$	1	1	$b^k$
<hr/>			
total	$\geq b^k$	$\leq b^k \left( \frac{b}{b-1} \right)^2$	

Numerical comparison for  $k = 5$  and  $b = 10$ :

$$\text{BFS} = 10 + 100 + 1,000 + 10,000 + 100,000 = 111,110$$

$$\text{IDS} = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$$

*Note:* IDS recalculates shallow nodes several times, but this doesn't have a big effect compared to BFS!

# BIDIRECTIONAL SEARCH (3.4.6)

## DIRECTION OF SEARCH

The definition of searching is symmetric: find path from start nodes to goal node or from goal node to start nodes.

*Forward branching factor*: number of arcs going out from a node.

*Backward branching factor*: number of arcs going into a node.

Search complexity is  $O(b^n)$ . Therefore, we should use forward search if forward branching factor is less than backward branching factor, and vice versa.

Note: when a graph is dynamically constructed, the backwards graph may not be available.



# BIDIRECTIONAL SEARCH

*Idea:* search backward from the goal and forward from the start simultaneously.

- This can result in an exponential saving, because  $2b^{k/2} \ll b^k$ .
- The main problem is making sure the frontiers meet.

One possible implementation:

- Use BFS to gradually search backwards from the goal, building a set of locations that will lead to the goal.
  - this can be done using *dynamic programming*
- Interleave this with forward heuristic search (e.g., A\*) that tries to find a path to these interesting locations.

## DYNAMIC PROGRAMMING

*Idea:* for statically stored graphs, build a table of the actual distance  $dist(n)$ , of the shortest path from node  $n$  to a goal.

- This can be built backwards from the goal:

$$dist(n) = \begin{array}{ll} \text{if } isGoal(n) & \text{then } 0 \\ \text{else } \min_{(n,m) \in G} & (|(n,m)| + dist(m)) \end{array}$$

The calculation of  $dist$  can be interleaved with a forward heuristic search.

# MEMORY-BOUNDED A\* (3.5.3)

The biggest problem with A\* is the space usage.

Can we make an iterative deepening version?

- IDA\*: use the  $f$  value as the cutoff cost
  - the cutoff is the smallest  $f$  value that exceeded the previous cutoff
  - often useful for problems with unit step costs
  - **problem:** with real-valued costs, it risks regenerating too many nodes
- RBFS: recursive best-first search
  - similar to DFS, but continues along a path until  $f(n) > limit$
  - $limit$  is the  $f$  value of the best *alternative path* from an ancestor
  - if  $f(n) > limit$ , recursion unwinds to alternative path
  - **problem:** regenerates too many nodes
- SMA\* and MA\*: (simplified) memory-bounded A\*
  - uses all available memory
  - when memory is full, it drops the worst leaf node from the frontier

**LOCAL SEARCH (R&N 4.1)**

**HILL CLIMBING (4.1.1–4.1.2)**

**POPULATION-BASED METHODS (4.1.3–4.1.4)**

# ITERATIVE BEST IMPROVEMENT

In many optimization problems, the path is irrelevant

- the goal state itself is the solution

Then the state space can be the set of “complete” configurations

- e.g., for 8-queens, a configuration can be any board with 8 queens (it is irrelevant in which order the queens are added)

In such cases, we can use *iterative improvement* algorithms; we keep a single “current” state, and try to improve it

- e.g., for 8-queens, we start with 8 queens on the board, and gradually move some queen to a better place

The goal would be to find an optimal configuration

- e.g., for 8-queens, where no queen is threatened


This takes constant space, and is suitable for online and offline search

## EXAMPLE: **n**-QUEENS

Put **n** queens on an  $n \times n$  board, in separate columns


Move a queen to reduce the number of conflicts;  
repeat until we cannot move any queen anymore

⇒ then we are at a local maximum, hopefully it is global too



This almost always solves **n**-queens problems  
almost instantaneously for very large **n** (e.g., **n** = 1 million)

## EXAMPLE: 8-QUEENS




Move a queen within its column, choose the minimum n:o of conflicts

- the best moves are marked above (conflict value: 12)
- after 5 steps we reach a local minimum (conflict value: 1)

## EXAMPLE: TRAVELLING SALESPERSON

Start with any complete tour, and perform pairwise exchanges



Variants of this approach get within 1% of optimal very quickly with thousands of cities






# HILL CLIMBING SEARCH (4.1.1–4.1.2)

Also called gradient/steepest ascent/descent,  
or greedy local search.

```
function HillClimbing(graph, initialState):  
    current := initialState  
    loop:  
        neighbor := a highest-valued successor of current  
        if neighbor.value ≤ current.value then return current  
        current := neighbor
```

# PROBLEMS WITH HILL CLIMBING

Local maxima — Ridges — Plateaux



## RANDOMIZED ALGORITHMS

Consider two methods to find a minimum value:

- Greedy ascent: start from some position, keep moving upwards, and report maximum value found
- Pick values at random, and report maximum value found

Which do you expect to work better to find a global maximum?

Can a mix work better?

## RANDOMIZED HILL CLIMBING


As well as upward steps we can allow for:

- *Random steps*: (sometimes) move to a random neighbor.
- *Random restart*: (sometimes) reassign random values to all variables.

Both variants can be combined!

# 1-DIMENSIONAL ILLUSTRATIVE EXAMPLE

Two 1-dimensional search spaces; you can step right or left:



Which method would most easily find the global maximum?

- random steps or random restarts?

What if we have hundreds or thousands of dimensions?

- ...where different dimensions have different structure?

# SIMULATED ANNEALING

Simulated annealing is an implementation of random steps:

```
function SimulatedAnnealing(problem, schedule):  
  current := problem.initialState  
  for t in 1, 2, ...:  
    T := schedule(t)  
    if T = 0 then return current  
    next := a randomly selected neighbor of current  
     $\Delta E := next.value - current.value$   
    if  $\Delta E > 0$  or with probability  $e^{\Delta E/T}$ :  
      current := next
```

$T$  is the “cooling temperature”, which decreases slowly towards 0

The cooling speed is decided by the *schedule*

# POPULATION-BASED METHODS (4.1.3–4.1.4)

## LOCAL BEAM SEARCH

*Idea:* maintain a population of  $k$  states in parallel, instead of one.

- At every stage, choose the  $k$  best out of all of the neighbors.
  - when  $k = 1$ , it is normal hill climbing search
  - when  $k = \infty$ , it is breadth-first search
- The value of  $k$  lets us limit space and parallelism.
- *Note:* this is not the same as  $k$  searches run in parallel!
- *Problem:* quite often, all  $k$  states end up on the same local hill.

# STOCHASTIC BEAM SEARCH

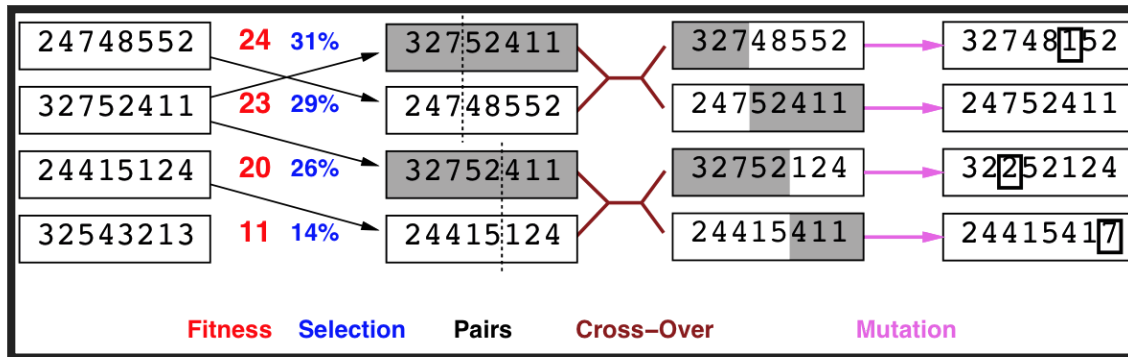
Similar to beam search, but it chooses the next  $k$  individuals *probabilistically*.

- The probability that a neighbor is chosen is proportional to its heuristic value.
- This maintains diversity amongst the individuals.
- The heuristic value reflects the fitness of the individual.
- Similar to natural selection:  
each individual mutates and the fittest ones survive.



# GENETIC ALGORITHMS

Similar to stochastic beam search,  
but *pairs* of individuals are combined to create the offspring.




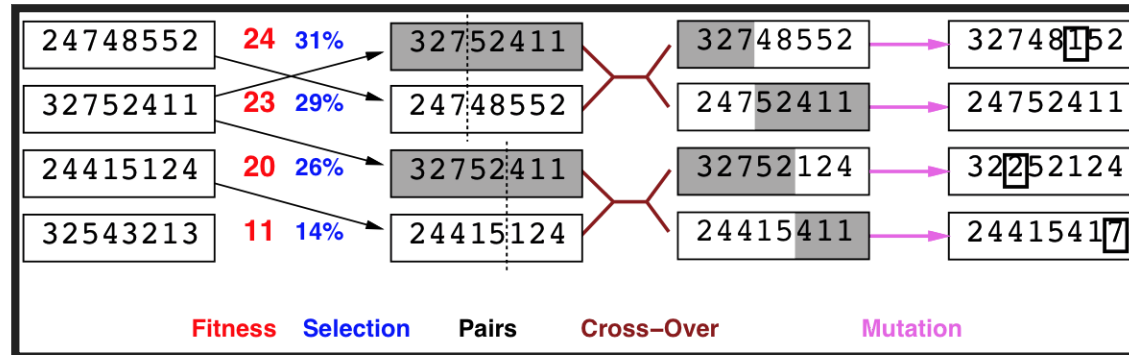
For each generation:

- Randomly choose pairs of individuals where the fittest individuals are more likely to be chosen.
- For each pair, perform a cross-over: form two offspring each taking different parts of their parents:
- Mutate some values.

Stop when a solution is found.

# $n$ QUEENS ENCODED AS A GENETIC ALGORITHM

The  $n$  queens problem can be encoded as  $n$  numbers  $1 \dots n$ :



# EVALUATING RANDOMIZED ALGORITHMS (NOT IN R&N)

How can you compare three algorithms A, B and C, when

- A solves the problem 30% of the time very quickly but doesn't halt for the other 70% of the cases
- B solves 60% of the cases reasonably quickly but doesn't solve the rest
- C solves the problem in 100% of the cases, but slowly?

Summary statistics, such as mean run time or median run time don't make much sense.

# RUNTIME DISTRIBUTION

Plots the runtime and the proportion of the runs that are solved within that runtime.

